

Class Segmentation and Object Localization with Superpixel Neighborhoods

Brian Fulkerson¹ Andrea Vedaldi² Stefano Soatto¹

¹ Department of Computer Science
University of California, Los Angeles, CA 90095
`{bfulkers, soatto}@cs.ucla.edu`

² Department of Engineering Science
University of Oxford, UK
`vedaldi@robots.ox.ac.uk`

Abstract

We propose a method to identify and localize object classes in images. Instead of operating at the pixel level, we advocate the use of superpixels as the basic unit of a class segmentation or pixel localization scheme. To this end, we construct a classifier on the histogram of local features found in each superpixel. We regularize this classifier by aggregating histograms in the neighborhood of each superpixel and then refine our results further by using the classifier in a conditional random field operating on the superpixel graph. Our proposed method exceeds the previously published state-of-the-art on two challenging datasets: Graz-02 and the PASCAL VOC 2007 Segmentation Challenge.

1. Introduction

Recent success in image-level object categorization has led to significant interest on the related fronts of localization and pixel-level categorization. Both areas have seen significant progress, through object detection challenges like PASCAL VOC [9]. So far, the most promising techniques seem to be those that consider each pixel of an image.

For localization, sliding window classifiers [8, 3, 21, 35] consider a window (or all possible windows) around each pixel of an image and attempt to find the classification which best fits the model. Lately, this model often includes some form of spatial consistency (e.g. [22]). In this way, we can view sliding window classification as a “top-down” localization technique which tries to fit a coarse global object model to each possible location.

In object class segmentation, the goal is to produce a pixel-level segmentation of the input image. Most approaches are built from the bottom up on learned local representations (e.g. TextronBoost [32]) and can be seen as an evolution of texture detectors. Because of their rather local nature, a conditional random field [20] or some other model is often introduced to enforce spatial consistency. For computational reasons, this usually operates on a re-

duced grid of the image, abandoning pixel accuracy in favor of speed. The current state-of-the-art for the PASCAL VOC 2007 Segmentation Challenge [31] is a scheme which falls into this category.

Rather than using the pixel grid, we advocate a representation adapted to the local structure of the image. We consider small regions obtained from a conservative oversegmentation, or “superpixels,” [29, 10, 25] to be the elementary unit of any detection, categorization or localization scheme.

On the surface, using superpixels as the elementary units seems counter-productive, because aggregating pixels into groups entails a decision that is unrelated to the final task. However, aggregating pixels into superpixels captures the local redundancy in the data, and the goal is to perform this decision in a conservative way to minimize the risk of merging unrelated pixels [33]. At the same time, moving to superpixels allows us to measure feature statistics (in this case: histograms of visual words) on a naturally adaptive domain rather than on a fixed window. Since superpixels tend to preserve boundaries, we also have the opportunity to create a very accurate segmentation by simply finding the superpixels which are part of the object.

We show that by aggregating neighborhoods of superpixels we can create a robust region classifier which exceeds the state-of-the-art on Graz-02 pixel-localization and on the PASCAL VOC 2007 Segmentation Challenge. Our results can be further refined by a simple conditional random field (CRF) which operates on superpixels, which we propose in Section 3.4.

2. Related Work

Sliding window classifiers have been well explored for the task of detecting the location of an object in an image [3, 21, 8, 9]. Most recently, Blaschko *et al.* [3] have shown that it is feasible to search all possible sub-windows of an image for an object using branch and bound and a structured classifier whose output is a bounding box. However, for our purposes a bounding box is not an acceptable final output, even for the task of localization.

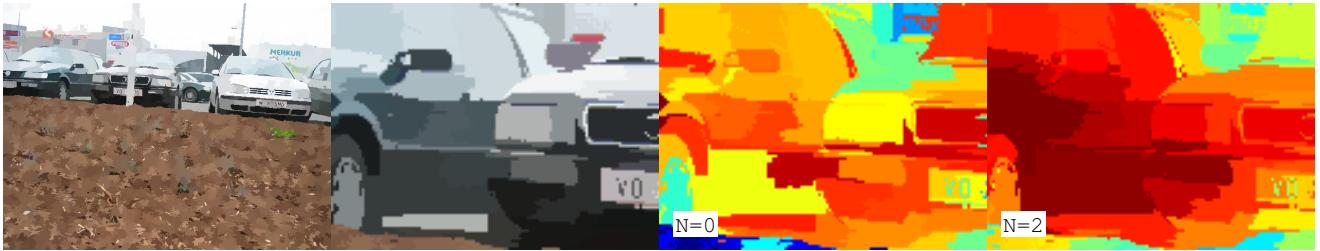


Figure 1. Aggregating histograms. An illustration of the detail of our superpixel segmentation and the effectiveness of aggregating histograms from adjacent segments. From the left: segmentation of a test image from Graz-02, a zoomed in portion of the segmentation, the classification of each segment where more red is more car-like, and the resulting classification after aggregating all histograms within $N = 2$ distance from the segment being classified.

Our localization capability is more comparable to Marszałek [24] or Fulkerson *et al.* [11]. Marszałek warps learned shape masks into an image based on distinctive local features. Fulkerson performs bag-of-features classification within a local region, as we do, but the size of the region is fixed (a rectangular window). In contrast, our method provides a natural neighborhood size, expressed in terms of low level image regions (the superpixels). A comparison with these methods is provided in Table 2.

Class segmentation algorithms which operate at the pixel level are often based on local features like textons [32] and are augmented by a conditional random field or another spatial coherency aid [15, 19, 16, 37, 17, 28, 13] to refine the results. Shotton *et al.* [31] constructs semantic texton forests for extremely fast classification. Semantic texton forests are essentially randomized forests of simple texture classifiers which are themselves randomized forests. We compare our results with and without an explicit spatial aid (a CRF) with those of Shotton in Table 3. Another notable work in this area is that of Gould *et al.* [13] who recently proposed a superpixel-based CRF which learns relative location offsets of categories. We eventually augment our model with a CRF on superpixels, but we do not model the relative locations of objects explicitly, instead preferring to use stronger local features and learn context via connectedness in the superpixel graph.

A number of works utilize one or more segmentations as a starting point for their task. An early example is Barnard *et al.* [2], who explore associating labels with image regions using simple color features and then merging regions based on similarity over the segment-label distribution. More recently, Russell *et al.* [30] build a bag-of-features representation on multiple segmentations to automatically discover object categories and label them in an unsupervised fashion. Similarly, Galleguillos *et al.* [12] use Multiple Instance Learning (MIL) to localize objects in weakly labeled data. Both assume that at least one of their segmentations contains a segment which correctly separates the entire object from the background. By operating

on superpixels directly, we can avoid this assumption and the associated difficulty of finding the one “good” segment.

Perhaps the most closely related work to ours is that of Pantofaru *et al.* [27]. Pantofaru *et al.* form superpixel-like objects by intersecting multiple segmentations and then classify these by averaging the classification results from all of the member regions. Their model allows them to gather classification information from a number of different neighborhoods sizes (since each member segment has a different extent around the region being classified). However, multiple segmentations are much more computationally expensive than superpixels, and we exceed their performance on the VOC 2007 dataset (see Table 3).

Additionally, a number of authors use graphs of image structures for various purposes, including image categorization [14, 26] and medical image classification [1]. Although we operate on a graph, we do not seek to mine discriminative substructures [26] or classify images based on the similarity of walks [14]. Instead we use the graph only to define neighborhoods and optionally to construct a conditional random field.

3. Superpixel Neighborhoods

3.1. Superpixels

We use quick shift [36] to extract superpixels from our input images. Our model is quite simple: we perform quick shift on a five-dimensional vector composed of the LUV colorspace representation of each pixel and its location in the image.

Unlike superpixelization schemes based on normalized cuts (e.g. [29]), the superpixels produced by quick shift are not fixed in approximate size or number. A complex image with many fine scale image structures may have many more superpixels than a simple one, and there is no parameter which puts a penalty on the boundary, leading to superpixels which are quite varied in size and shape. Statistics related to our superpixels (such as the average size and degree in the graph) are detailed in Section 4.

This produces segmentations, like the one in Figure 1, which consist of many small regions that preserve most of the boundaries in the original image. Since we perform this segmentation on the full resolution image, we leave open the potential to obtain a nearly pixel-perfect segmentation of the object.

3.2. Classification

We construct a bag-of-features classifier which operates on the regions defined by the superpixels we have found. SIFT descriptors [23] are extracted for each pixel of the image at a fixed scale and orientation using the fast SIFT framework found in [34]. The extracted descriptors are then quantized using a K -means dictionary and aggregated into one l^1 -normalized histogram $h_i^0 \in \mathbb{R}_+^K$ for each superpixel $s_i \in S$. In order to train the classifier, each superpixel s_i is assigned the most frequent class label it contains (since the ground-truth labels have pixel-level granularity). Then a one-vs-rest support vector machine (SVM) with an RBF- χ^2 kernel is trained on the labeled histograms for each of the object categories. This yields discriminant functions of the form

$$C(h^0) = \sum_{j=1}^L c_j \exp(-\gamma d_{\chi^2}^2(h^0, h_j^0))$$

where $c_j \in \mathbb{R}$ are coefficients and h_j^0 representative histograms (support vectors) selected by SVM training, $\gamma \in \mathbb{R}_+$ is a parameter selected by cross-validation, and $d_{\chi^2}^2(h^0, h_j^0)$ is the χ^2 distance between histograms h^0 and h_j^0 , defined as

$$d_{\chi^2}^2(h^0, h_j^0) = \sum_{k=1}^K \frac{(h^0(k) - h_j^0(k))^2}{h^0(k) + h_j^0(k)}.$$

The classifier which results from this is very specific. It finds superpixels which resemble superpixels that were seen in the training data without considering the surrounding region. This means that while a wheel or grille on a car may be correctly identified, the nearby hub of the wheel or the headlight can be detected with lower confidence or missed altogether (Figure 1).

Another drawback of learning a classifier for each superpixel is that the histograms associated with each superpixel are very sparse, often containing only a handful of non-zero elements. This is due to the nature of our superpixels: by definition they cover areas that are roughly uniform in color and texture. Since our features are fixed-scale and extracted densely, our superpixels sometimes contain tens or even hundreds of descriptors that quantize to the same visual word.

3.3. Superpixel Neighborhoods

We address both of the problems mentioned in the previous section by introducing histograms based on superpixel neighborhoods. Let $G(S, E)$ be the adjacency graph of superpixels $s_i \in S$ in an image, and H_i^0 be the unnormalized histogram associated with this region. E is the set of edges formed between pairs of adjacent superpixels (s_i, s_j) in the image and $D(s_i, s_j)$ is the length of shortest path between two superpixels. Then, H_i^N is the histogram obtained by merging the histograms of the superpixel s_i and neighbors who are less than N nodes away in the graph:

$$H_i^N = \sum_{s_j | D(s_i, s_j) \leq N} H_j^0$$

The learning framework is unchanged, except that we describe superpixels by the histograms $h_i^N = H_i^N / \|H_i^N\|_1$ in place of h_i^0 .

Using these histograms in classification addresses both of our previous issues. First, since adjacent superpixels must be visually dissimilar, histograms constructed from superpixel neighborhoods contain more diverse features and are therefore less sparse. This provides a regularization for our SVM, reducing overfitting. It also provides spatial consistency in our classification because as we increase N , histograms of adjacent superpixels have more features in common.

Second, because we are effectively increasing the spatial extent of the region considered in classification, we are also providing our classifier with a better description of the object. As we increase N we move from the “part” level to the “object” level, and since not all training superpixels will lie on the interior of the object, we are also learning some “context”.

However, note that as N becomes larger we will blur the boundaries of our objects since superpixels which are on both sides of the object boundary will have similar histograms. In the next section, we explore adding a CRF to reduce this effect.

3.4. Refinement with a CRF

In order to recover more precise boundaries while still maintaining the benefits of increasing N , we must introduce new constraints that allow us to reduce misclassifications that occur near the edges of objects. Conditional random fields provide a natural way to incorporate such constraints by including them in the pairwise edge potential of the model. Let $P(\mathbf{c}|G; w)$ be the conditional probability of the set of class label assignments \mathbf{c} given the adjacency graph $G(S, E)$ and a weight w :

$$-\log(P(\mathbf{c}|G; w)) = \sum_{s_i \in S} \Psi(c_i | s_i) + w \sum_{(s_i, s_j) \in E} \Phi(c_i, c_j | s_i, s_j)$$

Our unary potentials Ψ are defined directly by the probability outputs provided by our SVM [7] for each superpixel:

$$\Psi(c_i|s_i) = -\log(P(c_i|s_i))$$

and our pairwise edge potentials Φ are similar to those of [32, 6]:

$$\Phi(c_i, c_j|s_i, s_j) = \left(\frac{L(s_i, s_j)}{1 + \|s_i - s_j\|} \right) [c_i \neq c_j]$$

where $[\cdot]$ is the zero-one indicator function and $\|s_i - s_j\|$ is the norm of the color difference between superpixels in the LUV colorspace. $L(s_i, s_j)$ is the shared boundary length between superpixels s_i and s_j and acts here as a regularizing term which discourages small isolated regions.

In many CRF applications for this domain, the unary and pairwise potentials are represented by a weighted summation of many simple features (e.g. [32]), and so the parameters of the model are learned by maximizing their conditional log-likelihood. In our formulation, we simply have one weight w which represents the tradeoff between spatial regularization and our confidence in the classification. We estimate w by cross validation on the training data. Once our model has been learned, we carry out inference with the multi-label graph optimization library of [4, 18, 5] using α -expansion. Since the CRF is defined on the superpixel graph, inference is very efficient, taking less than half a second per image.

Results with the CRF are presented in Section 4 as well as Figures 2 and 3.

4. Experiments

We evaluate our algorithm for varying N with and without a CRF on two challenging datasets. Graz-02 contains three categories (bicycles, cars and people) and a background class. The task is to localize each category against the background class. Performance on this dataset is measured by the pixel precision-recall.

The PASCAL VOC 2007 Segmentation Challenge [9] contains 21 categories and few training examples. While the challenge specifies that the detection challenge training data may also be used, we use only the ground truth segmentation data for training. The performance measure for this dataset is the average pixel accuracy: for each category the number of correctly classified pixels is divided by the ground truth pixels plus the number of incorrectly classified pixels. We also report the total percentage of pixels correctly classified.

MATLAB code to reproduce our experiments is available from our website¹.

¹<http://vision.ucla.edu/bag/>

4.1. Common Parameters

Experiments on both datasets share many of the same parameters which we detail here.

SIFT descriptors are extracted at each pixel with a patch size of 12 pixels and fixed orientation. These descriptors are quantized into a K -means dictionary learned on the training data. All experiments we present here use $K = 400$, though in Figure 1 we show that a wide variety of K produce similar results.

The superpixels extracted via quick shift are controlled by three parameters: λ , the trade-off between color importance and spatial importance, σ , the scale at which the density is estimated, and τ , the maximum distance in the feature space between members of the same region. We use the same parameters for all of our experiments: $\sigma = 2$, $\lambda = 0.5$, $\tau = 8$. These values were determined by segmenting a few training images from Graz-02 by hand until we found a set which preserved nearly all of the object boundaries and had the largest possible average segment size. In principle, we could do this search automatically on the training data, looking for the parameter set which creates the largest average segment size while ensuring that the maximum possible classification accuracy is greater than some desired level. In practice, the algorithm is not too sensitive to the choice of parameters, so a quick tuning by hand is sufficient. Note that the number or size of the superpixels is not fixed (as opposed to [13]): the selected parameters put a rough bound on the maximum size of the superpixels but do not control the shape of the superpixels or degree of the superpixel graph.

Histograms for varying N are extracted as described in Section 3.3 and labels are assigned to training superpixels by the majority class vote. We randomly select an equal number of training histograms from each category as the training data for our SVM.

We learn a one-vs-rest multi-class SVM with an RBF- χ^2 kernel on the histograms using libsvm [7] as described in Section 3.2. During testing, we convert our superpixel labels into a pixel-labeled map and evaluate at the pixel level for direct comparison with other methods.

In both experiments, we take our final SVM and include it in the CRF model described in Section 3.4.

4.2. Graz-02

On Graz-02, we use the same training and testing split as Marszałek and Schmid [24] and Fulkerson *et al.* [11]. Our segment classifier is trained on 750 segments collected at random from the category and the background.

Graz-02 images are 640 by 480 pixels and quick shift produces approximately 2000 superpixels per image with an average size of 150 pixels. The average degree of the su-

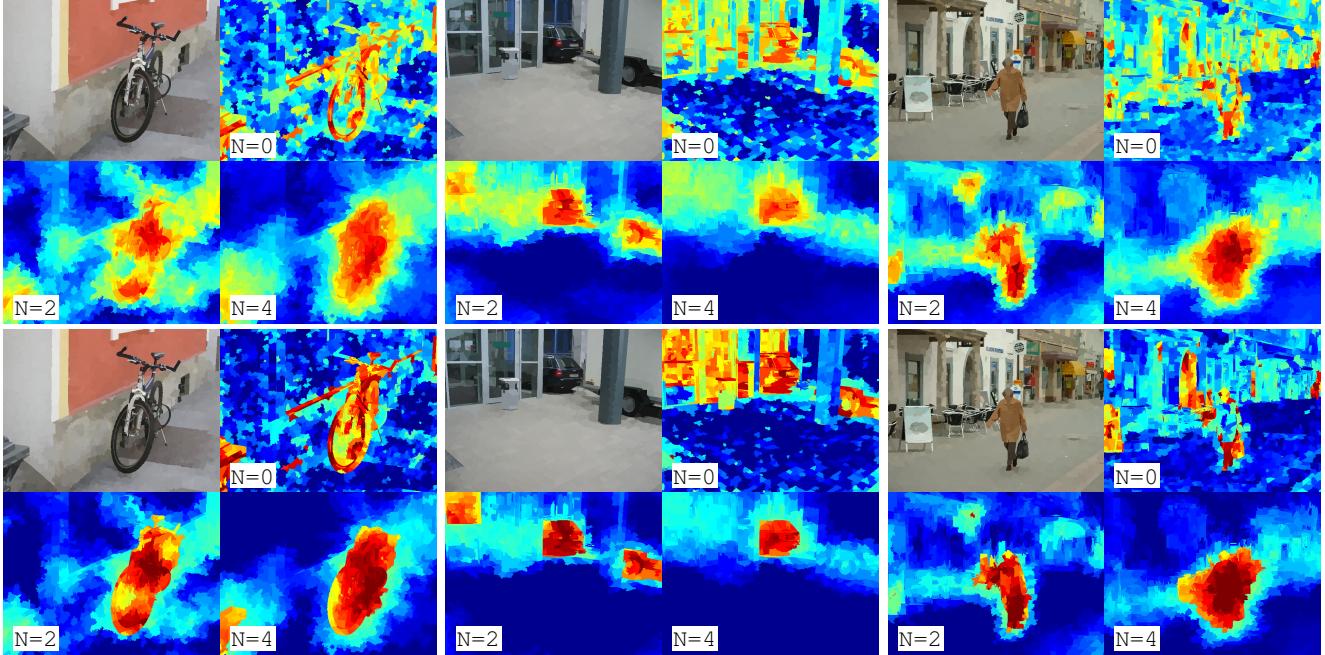


Figure 2. Graz-02 confidence maps. Our method produces very well localized segmentations of the target category on Graz-02. Here, a dark red classification means that the classifier is extremely confident the region is foreground (using the probability output of libsvm), while a dark blue classification indicates confident background classification. Notice that as we increase the number of neighbors considered (N) regions which were uncertain become more confident and spurious detections are suppressed. **Top two rows:** Without CRF. **Bottom two rows:** With CRF.

	Graz-02 $N =$					PASCAL 2007 $N =$				
	0	1	2	3	4	0	1	2	3	4
$K = 10$	37	44	47	51	49	10	10	12	12	12
$K = 100$	48	61	64	64	64	13	19	23	25	25
$K = 200$	49	63	66	66	64	13	20	25	26	25
$K = 400$	50	64	67	69	67	14	21	25	28	27
$K = 1000$	49	63	68	68	66	14	22	27	27	26

Table 1. Effect of K . Here we explore the effect of the dictionary size K on the accuracy of our method (without a CRF) for varying neighborhood sizes N . Increasing the size of the dictionary increases performance until we begin to overfit the data. We pick $K = 400$ for our experiments, but a large range of K will work well. Notice that even with $K = 10$ we capture some information, and increasing N still provides noticeable improvement.

perpixel graph is 6, however the maximum degree is much larger (137).

In Table 2, we compare our results for varying size N with those of Fulkerson *et al.* [11] which uses a similar bag-of-features framework and Marszalek and Schmid [24] which warps shape masks around likely features to define probable regions. We improve upon the state-of-the-art in all categories (+17% on cars, +15% on people, and +6% on bicycles).

Example localizations may be found in Figures 1 and 2.

	Cars	People	Bicycles
[24] full framework	53.8%	44.1%	61.8%
[11] NN	54.7%	47.1%	66.4%
[11] SVM	49.4%	51.4%	65.2%
$N = 0$	43.3%	51.3%	56.7%
CRF $N = 0$	46.0%	54.3%	63.4%
$N = 1$	62.0%	62.7%	67.6%
CRF $N = 1$	69.7%	63.8%	69.7%
$N = 2$	67.1%	65.4%	69.3%
CRF $N = 2$	71.2%	66.3%	71.2%
$N = 3$	68.6%	65.7%	71.7%
CRF $N = 3$	72.2%	66.1%	72.2%
$N = 4$	67.1%	62.7%	71.0%
CRF $N = 4$	71.3%	63.2%	71.3%

Table 2. Graz-02 results. The precision = recall points for our experiments on Graz-02. Compared to the former state-of-the-art [11], we show a 17% improvement on Cars, a 15% improvement on People and a 6% improvement on Bicycles. N is the distance of the furthest neighboring region to aggregate, as described in Section 3.3. Our best performing case is always the CRF-augmented model described in Section 3.4.

Notice that although $N = 0$ produces some very precisely defined correct classifications, there are also many missed

detections and false positives. As we increase the amount of local information that is considered for each classification, regions that were classified with lower confidence become more confident, and false positives are suppressed.

Adding the CRF provides consistent improvement, sharpening the boundaries of objects and providing further spatial regularization. Our best performing cases use $N = 2$ or $N = 3$, balancing the incorporation of extra local support with the preference for compact regions with regular boundaries.

4.3. VOC 2007 Segmentation

For the VOC challenge, we use the same sized dictionary and features as Graz-02 ($K = 400$, patch size = 12 pixels). The training and testing split is defined in the challenge. We train on the training and validation sets and test on the test set. Since there are fewer training images per category, for this experiment we train on 250 randomly selected training histograms from each category.

VOC 2007 images are not fixed size and tend to be smaller than those in Graz-02, so with the same parameters quick shift produces approximately 1200 superpixels per image with a mean size of 150 pixels. The average degree of the superpixel graph is 6.4, and the maximum degree is 72.

In Table 3 we compare with the only segmentation entry in the challenge (Oxford Brookes), as well as the results of Shotton *et al.* [31], and Pantofaru *et al.* [27]. Note that Shotton reports a set of results which bootstrap a detection entry (TKK). We do not compare with these results because we do not have the data to do so. However, because our classifier is simply a multi-class SVM, we can easily add either the Image Level Prior (ILP) or a Detection Level Prior (DLP) that Shotton uses. Even without the ILP, we find that we outperform Shotton with the ILP on 14 of the 21 categories and tie on one more. Our average performance is also improved by 8%. Compared to Shotton without ILP or Pantofaru, average performance is improved by 12%. Selected segmentations may be found in Figure 3.

This dataset is much more challenging (we are separating 21 categories instead of 2, with less training data and more variability) and because of this when $N = 0$ everything has very low confidence. As we increase N we start to see contextual relationships playing a role. For example, in the upper left image of Figure 4 we see that as the person classification gets more confident, so does the bike and motorbike classification, since this configuration (person above bike) occurs often in the training data. We also see that larger N tends to favor more contiguous regions, which is consistent with what we expect to observe.

On this dataset, adding a CRF improves the qualitative results significantly, and provides a consistent boost for the accuracy as well. Object boundaries become crisp, and of-

ten the whole object has the same label, even if it is not always the correct one.

5. Conclusion

We have demonstrated a method for localizing objects and segmenting object classes that considers the image at the level of superpixels. Our method exceeds the state-of-the-art on Graz-02 and the PASCAL VOC 2007 Segmentation Challenge, even without the aid of a CRF or color information. When we add a CRF which penalizes pairs of superpixels that are very different in color, we consistently improve both our quantitative and especially our qualitative results.

Acknowledgments

This research was supported by ONR 67F-1080868/N0014-08-1-0414 and NSF 0622245.

References

- [1] E. Aldea, J. Atif, and I. Bloch. Image classification using marginalized kernels for graphs. In *Proc. CVPR*, 2007.
- [2] K. Barnard, P. Duygulu, R. Guru, P. Gabbur, and D. Forsyth. The effects of segmentation and feature choice in a translation model of object recognition. In *Proc. CVPR*, 2003.
- [3] M. Blaschko and C. Lampert. Learning to localize objects with structured output regression. In *Proc. ECCV*, 2008.
- [4] Y. Boykov and V. Kolmogorov. An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision. In *PAMI*, 2004.
- [5] Y. Boykov, O. Veksler, and R. Zabih. Efficient approximate energy minimization via graph cuts. In *PAMI*, 2001.
- [6] Y. Y. Boykov and M.-P. Jolly. Interactive graph cuts for optimality boundary & region segmentation of objects in N-D images. In *Proc. ICCV*, 2001.
- [7] C.-C. Chang and C.-J. Lin. *LIBSVM: a library for support vector machines*, 2001. Software available at <http://www.csie.ntu.edu.tw/~cjlin/libsvm>.
- [8] N. Dalal and B. Triggs. Histograms of oriented gradients for human detection. In *Proc. CVPR*, 2005.
- [9] M. Everingham, L. Van Gool, C. K. I. Williams, J. Winn, and A. Zisserman. The PASCAL Visual Object Classes Challenge 2007 (VOC2007) Results. <http://www.pascal-network.org/challenges/VOC/voc2007/workshop/index.html>.
- [10] P. F. Felzenszwalb and D. P. Huttenlocher. Efficient graph-based image segmentation. *IJCV*, 59(2), 2004.
- [11] B. Fulkerson, A. Vedaldi, and S. Soatto. Localizing objects with smart dictionaries. In *Proc. ECCV*, 2008.
- [12] C. Galleguillos, B. Babenko, A. Rabinovich, and S. Belongie. Weakly supervised object localization with stable segmentations. In *Proc. ECCV*, 2008.
- [13] S. Gould, J. Rodgers, D. Cohen, G. Elidan, and D. Koller. Multi-class segmentation with relative location prior. In *IJCV*, 2008.

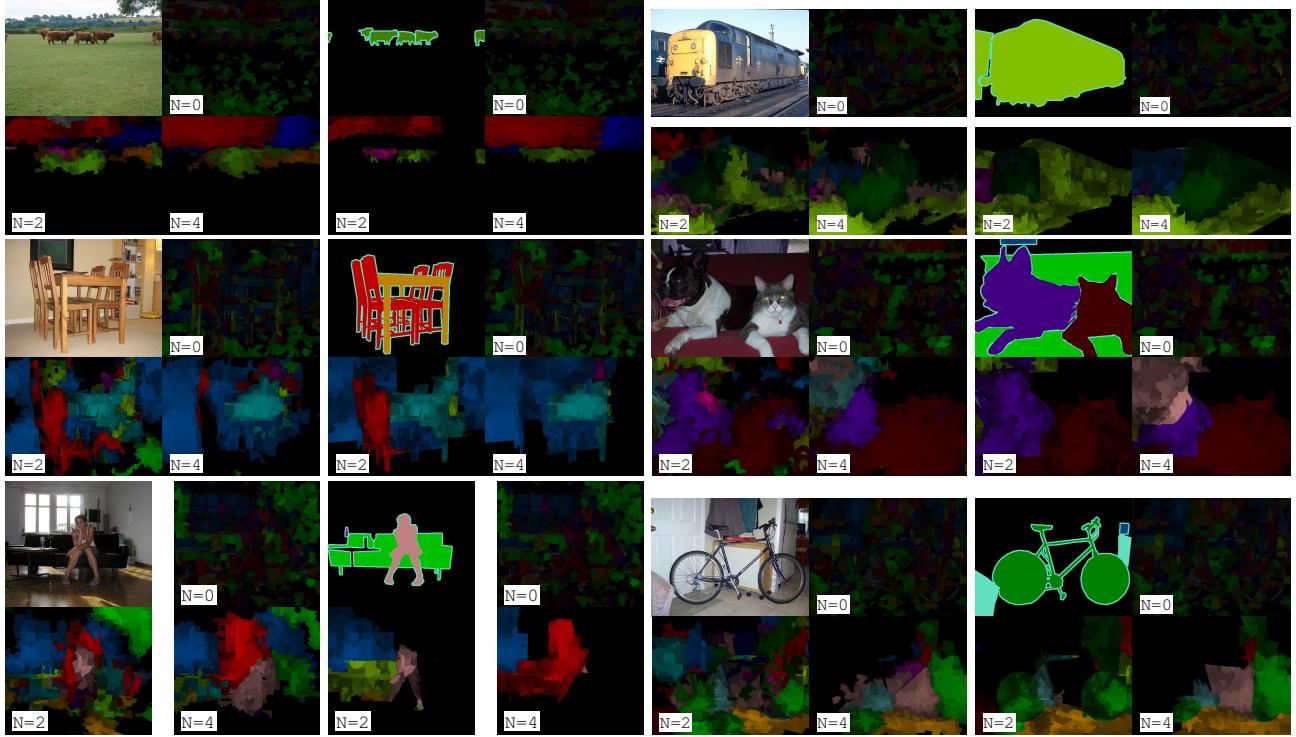


Figure 3. **PASCAL VOC 2007 + CRF.** Some selected segmentations for PASCAL. For each test image, the results are arranged into two blocks of four images. The first block (left-to-right) shows the results of the superpixel neighborhoods without a CRF. The second block uses the CRF described in Section 3.4. Colors indicate category and the intensity of the color is proportional to the posterior probability of the classification. Best viewed in color.

	background	aeroplane	bicycle	bird	boat	bottle	bus	car	cat	chair	cow	diningtable	dog	horse	motorbike	pottedplant	sheep	sofa	train	tvmonitor	Average	% Pixels	
Brookes	78	6	0	0	0	0	9	5	10	1	2	11	0	6	6	29	2	2	0	11	1	9	-
[27]	59	27	1	8	2	1	32	14	14	4	8	32	9	24	15	81	11	26	1	28	17	20	-
[31]	33	46	5	14	11	14	34	8	6	3	10	39	40	28	23	32	19	19	8	24	9	20	-
[31] + ILP	20	66	6	15	6	15	32	19	7	7	13	44	31	44	27	39	35	12	7	39	23	24	-
<i>N = 0</i>	21	14	8	8	17	14	10	7	19	13	13	7	16	9	13	2	10	23	34	17	20	14	18
CRF+ <i>N = 0</i>	20	14	8	8	17	14	10	7	19	13	13	7	16	9	13	2	10	23	34	17	20	14	18
<i>N = 1</i>	27	27	20	17	14	12	18	11	37	18	7	14	26	19	35	18	13	21	25	31	25	21	25
CRF+ <i>N = 1</i>	38	32	20	13	17	10	20	11	52	17	7	14	31	21	39	28	14	12	28	42	33	24	34
<i>N = 2</i>	36	27	26	15	11	5	26	29	42	25	9	15	36	23	58	32	17	11	20	37	29	25	34
CRF+ <i>N = 2</i>	56	26	29	19	16	3	42	44	56	23	6	11	62	16	68	46	16	10	21	52	40	32	51
<i>N = 3</i>	47	22	24	17	11	6	35	25	46	19	8	19	33	29	62	47	16	20	26	37	29	28	43
CRF+ <i>N = 3</i>	65	22	28	32	2	4	40	30	61	10	3	20	35	24	72	62	16	23	20	44	30	30	57
<i>N = 4</i>	51	20	22	18	7	2	39	25	49	15	6	14	36	28	64	56	15	17	21	40	23	27	46
CRF+ <i>N = 4</i>	65	20	30	22	2	2	39	25	57	10	3	7	36	23	66	62	15	17	8	46	11	27	57

Table 3. **VOC 2007 segmentation results.** Our best overall average performance (CRF+*N* = 2) performs better than Shotton *et al.* [31] with or without an Image Level Prior (ILP) on 14 out of 21 categories. Note that we could add ILP to our model. Similarly, we do not compare with the Shotton *et al.* results which used TKK’s detection results as a Detection Level Prior (DLP) because TKK’s detections were not available. We expect our method would provide a similar performance boost with this information. The CRF provides consistent improvement in average accuracy and in the percentage of pixels which were correctly classified.

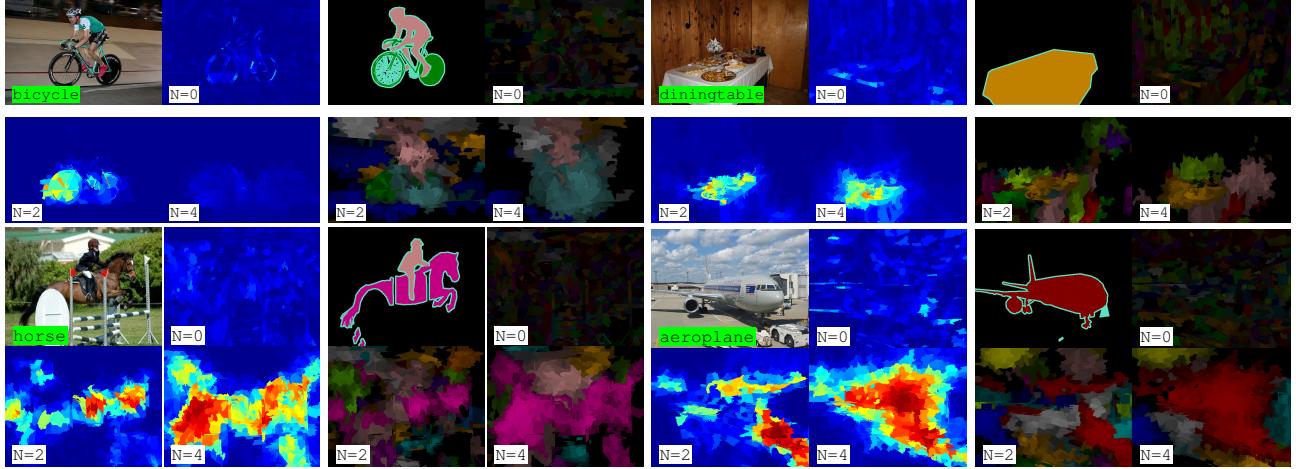


Figure 4. PASCAL VOC 2007 Confidence. Confidence maps for PASCAL. The results are arranged into two blocks of four images for each test image. The first block contains the input image, a category label, and the confidence map for that category for $N = 0, 2, 4$. The second block contains the ground truth labeling and our labellings with an intensity proportional to the confidence of the classification. Colors indicate category. For example, in the upper left we show the confidence for bicycle, and the classification which contains mostly bicycle (green) and some motorbike (light blue).

- [14] Z. Harchaoui and F. Bach. Image classification with segmentation graph kernels. In *Proc. CVPR*, 2007.
- [15] X. He, R. Zemel, and M. C.-P. nán. Multiscale conditional random fields for image labeling. In *Proc. CVPR*, 2004.
- [16] X. He, R. Zemel, and D. Ray. Learning and incorporating top-down cues in image segmentation. In *Proc. ECCV*, 2006.
- [17] D. Hoiem, C. Rother, and J. Winn. 3d layoutcrf for multi-view object class recognition and segmentation. In *Proc. CVPR*, 2007.
- [18] V. Kolmogorov and R. Zabih. What energy functions can be minimized via graph cuts? In *PAMI*, 2004.
- [19] S. Kumar and M. Hebert. A hierarchical field framework for unified context-based classification. In *Proc. ICCV*, 2005.
- [20] J. Lafferty, A. McCallum, and F. Pereira. Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *Proc. ICML*, 2001.
- [21] C. Lampert, M. Blaschko, and T. Hofmann. Beyond sliding windows: Object localization by efficient subwindow search. In *Proc. CVPR*, 2008.
- [22] S. Lazebnik, C. Schmid, and J. Ponce. Beyond bag of features: Spatial pyramid matching for recognizing natural scene categories. In *Proc. CVPR*, 2006.
- [23] D. G. Lowe. Distinctive image features from scale-invariant keypoints. *IJCV*, 2(60):91–110, 2004.
- [24] M. Marszałek and C. Schmid. Accurate object localization with shape masks. In *Proc. CVPR*, 2007.
- [25] A. Moore, S. Prince, J. Warrell, U. Mohammed, and G. Jones. Superpixel lattices. In *Proc. CVPR*, 2008.
- [26] S. Nowozin, K. Tsuda, T. Uno, T. Kudo, and G. BakIr. Weighted substructure mining for image analysis. In *Proc. CVPR*, 2007.
- [27] C. Pantofaru, C. Schmid, and M. Hebert. Object recognition by integrating multiple image segmentations. In *Proc. ECCV*, 2008.
- [28] A. Rabinovich, A. Vedaldi, C. Galleguillos, E. Wiewiora, and S. Belongie. Objects in context. In *Proc. ICCV*, 2007.
- [29] X. Ren and J. Malik. Learning a classification model for segmentation. In *Proc. ICCV*, 2003.
- [30] B. C. Russell, A. A. Efros, J. Sivic, W. T. Freeman, and A. Zisserman. Using multiple segmentations to discover objects and their extent in image collections. In *Proc. CVPR*, 2006.
- [31] J. Shotton, M. Johnson, and R. Cipolla. Semantic texton forests for image categorization and segmentation. In *Proc. CVPR*, 2008.
- [32] J. Shotton, J. Winn, C. Rother, and A. Criminisi. *Tenton-Boost*: Joint appearance, shape and context modeling for multi-class object recognition and segmentation. In *Proc. ECCV*, 2006.
- [33] S. Soatto. Actionable information in vision. In *Proc. ICCV*, October 2009.
- [34] A. Vedaldi and B. Fulkerson. VLFeat: An open and portable library of computer vision algorithms. <http://www.vlfeat.org/>, 2008.
- [35] A. Vedaldi, V. Gulshan, M. Varma, and A. Zisserman. Multiple kernels for object detection. In *Proc. ICCV*, 2009.
- [36] A. Vedaldi and S. Soatto. Quick shift and kernel methods for mode seeking. In *Proc. ECCV*, 2008.
- [37] J. Verbeek and B. Triggs. Region classification with markov field aspect models. In *Proc. CVPR*, 2007.